Parallel Vectorize

The `parallel_vectorize.py` module contains a set of llvmpy code generators for creating mulithreaded *ufunc*. It depends on the new `numpy.fromfunc` for turning arbitrary function pointers into *ufunc*.

From LLVM Function

The `parallel_vectorize_from_func` method generates multithreaded *ufunc* from LLVM functions.

First, we will implement a workload function:

```
In [1]: from llvm_cbuilder import *
        from llvm_cbuilder import shortnames as C
        from llvm.core import *

        # Implement a workload
        class Square(CDefinition):
            _name_ = 'square'
            _retty_ = C.double              # 1 output: double
            _argtys_ = [('x', C.double)]    # 1 input: double

            def body(self, x):
                self.ret(x * x)

        m = Module.new('my_module')
        llvm_square = Square()(m)  # Generate a llvm function
        print(llvm_square)
```

```
define double @square(double %x) {
decl:
  %x1 = alloca double
  br label %body

body:                                                ; preds = %decl
  store double %x, double* %x1
  %0 = load double* %x1
  %1 = load double* %x1
  %2 = fmul double %0, %1
  ret double %2
}
```

Then, we will generate a *ufunc* from `llvm_square`:

```
In [2]: from llvm.ee import *
        engine = EngineBuilder.new(m).create()           # Generate JIT engine

        from numbapro.parallel_vectorize import parallel_vectorize_from_func
        ufunc_square = parallel_vectorize_from_func(llvm_square, engine)  # Generate UFunc
```

```
[139838658969616L] [[12, 12]] 1 1 [None]
```

We are ready to use `ufunc_square` as a regular *ufunc*.

```
In [3]: import numpy as np
        A = np.arange(10., dtype=np.double)
        ufunc_square(A)
```

```
Out[3]: array([  0.,   1.,   4.,   9.,  16.,  25.,  36.,  49.,  64.,  81.])
```

Here's another example that uses three inputs:

```
In [4]: class SumOfThree(CDefinition):
            _name_ = 'sum.of.three'
            _retty_ = C.int
            _argtys_ = [('x', C.int),
                        ('y', C.int),
                        ('z', C.int)]
            def body(self, x, y, z):
                self.ret( x + y + z )

        llvm_sum3 = SumOfThree()(m)
        ufunc_sum3 = parallel_vectorize_from_func(llvm_sum3, engine)
        A = np.arange(10, dtype=np.int32)
        B = A * 10
        C = A * 100
        ufunc_sum3(A, B, C)
```

```
        [139838658970848L] [[5, 5, 5, 5]] 3 1 [None]
Out[4]: array([  0, 111, 222, 333, 444, 555, 666, 777, 888, 999], dtype=int32)
```

Internals

There are four functions behind each multithreaded *ufunc*.

  I. the workload function (user defined);
 II. the thread worker function (`UFuncCoreGeneric`);
III. the thread manager function (`ParallelUFuncPlatform`);
IV. the ufunc entry point function (`SpecializedParallelUFunc`).

UFuncCoreGeneric specializes to a llvm function type. **It currently understands simple builtin scalar types (integers, float, double) only as arguments and return-type for the workload function.** It sends work items to the workload function and performs work-stealing when it has finished its own workqueue. Work-stealing uses atomic compare-exchange (or CAS) instruction to acquire ownership of a workqueue. Work-stealing is implemented in the `UFuncCore._do_work_stealing`. It can be disabled on platform that does not support atomic operations.

ParallelUFuncPlatform specializes to the maximum number of threads. It divides all works equally among all threads. Each thread executes the function generated by `UFuncCoreGeneric` once.

SpecializedParallelUFunc is the specialized *ufunc* entry point for a specific combination of workload, UFuncCoreGeneric and ParallelUFuncPlatform.

Here's an example that uses `SpecializedParallelUFunc` directly for the `SumOfThree` workload.

```
In [5]: import numbapro.parallel_vectorize as pv
        # specialize
        def_spuf = pv.SpecializedParallelUFunc(pv.ParallelUFuncPlatform(num_thread=2),
                                               pv.UFuncCoreGeneric(llvm_sum3.type.pointee),
                                               CFuncRef(llvm_sum3))
        # define
        spuf = def_spuf(m)
        print(spuf.name)
```

```
        specialized_parallel_ufunc_2_ufunc_worker.i32.i32.i32.i32_sum.of.three
```

`CFuncRef` also accepts arbitrary function pointer as long as the function type is provided.

```
In [6]:  # specialize
         fnty = llvm_sum3.type.pointee
         sum3ptr = engine.get_pointer_to_function(llvm_sum3)
         print("as function pointer: %x" % sum3ptr)
         def_spuf = pv.SpecializedParallelUFunc(pv.ParallelUFuncPlatform(num_thread=2),
                                                pv.UFuncCoreGeneric(fnty),
                                                CFuncRef('sum3.as.ptr', fnty, sum3ptr)) # name, type, ptr
         # define
         spuf = def_spuf(m)
         print(spuf.name)
```

```
as function pointer: 7f2eb9980740
specialized_parallel_ufunc_2_ufunc_worker.i32.i32.i32.i32_sum3.as.ptr
```